

# Convolutional Feature Maps

Elements of efficient (and accurate)  
CNN-based object detection

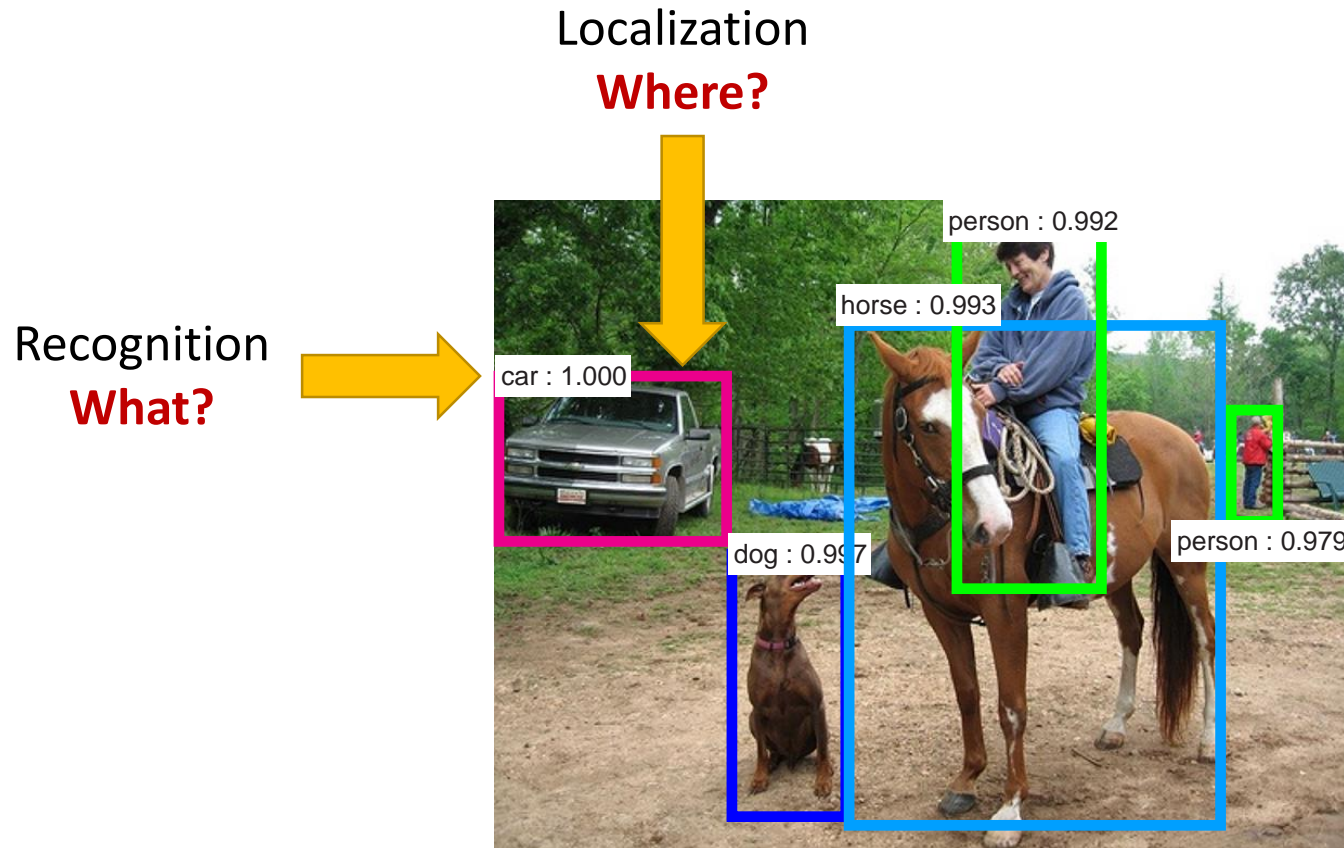
Kaiming He

Microsoft Research Asia (MSRA)

# Overview of this section

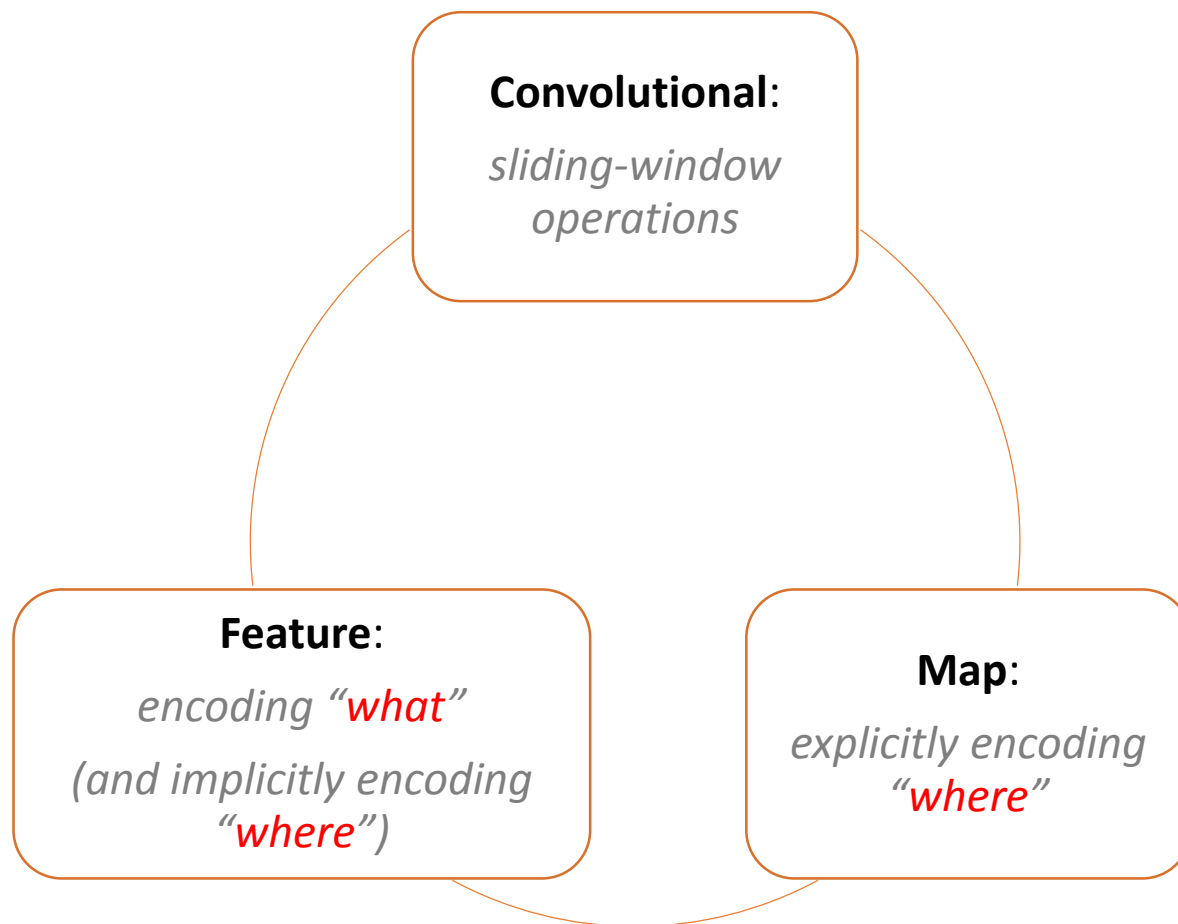
- Quick introduction to convolutional feature maps
  - Intuitions: into the “black boxes”
  - How object detection networks & region proposal networks are designed
  - Bridging the gap between “hand-engineered” and deep learning systems
- Focusing on forward propagation (inference)
  - Backward propagation (training) covered by Ross’s section

# Object Detection = What, and Where




- We need a building block that tells us “what and where”...

# Object Detection = What, and Where



# Convolutional Layers

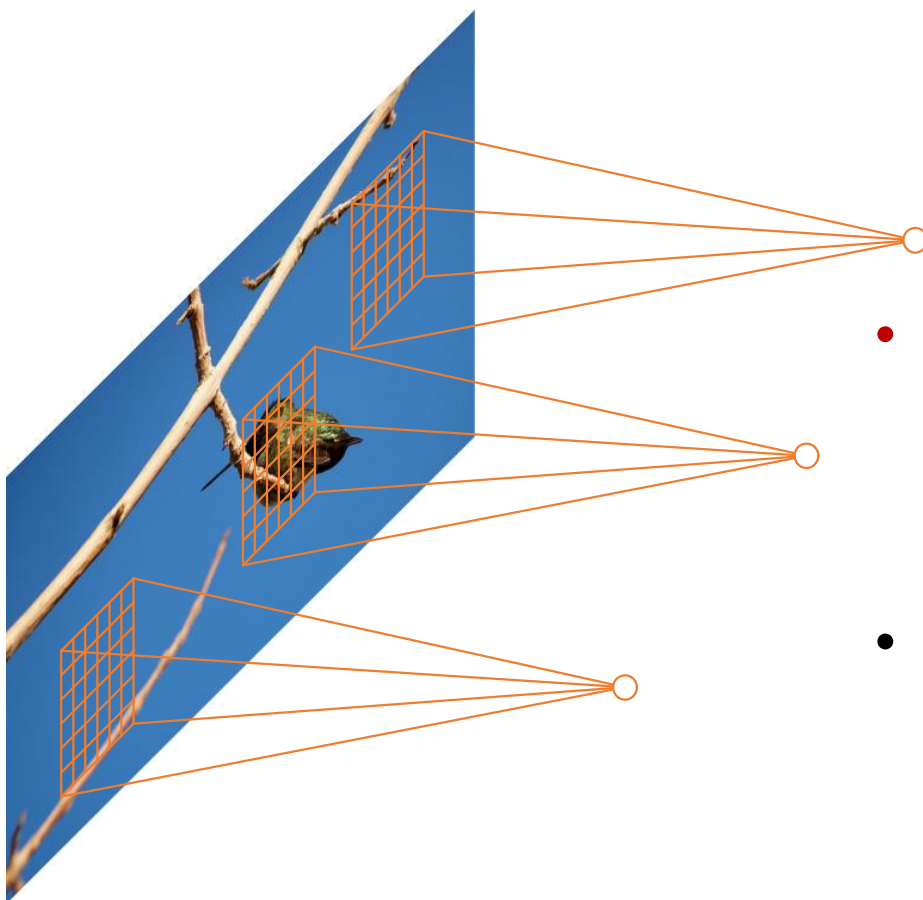
- Convolutional layers are **locally connected**



- a filter/kernel/window slides on the image or the previous map
- the **position** of the filter explicitly provides information for localizing
- local spatial information w.r.t. the window is encoded in the channels

# Convolutional Layers


- Convolutional layers share weights spatially: **translation-invariant**



- **Translation-invariant**: a translated region will produce the same response at the correspondingly translated position
- A local pattern's convolutional response can be **re-used** by different candidate regions

# Convolutional Layers

- Convolutional layers can be applied to **images of any sizes**, yielding **proportionally-sized** outputs



# HOG by Convolutional Layers

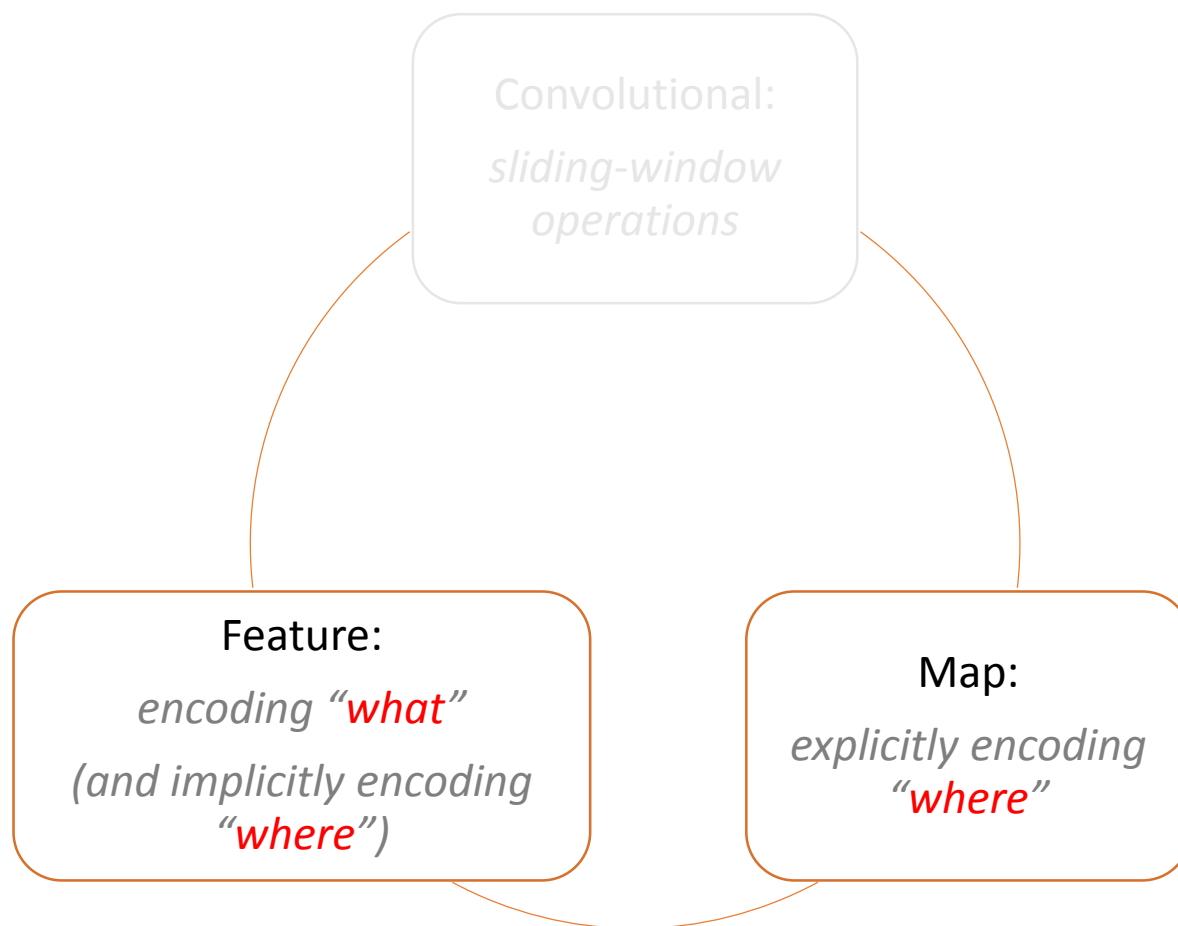
- Steps of computing HOG:
  - Computing image gradients
  - Binning gradients into 18 directions
  - Computing cell histograms
  - Normalizing cell histograms
- Convolutional perspectives:
  - Horizontal/vertical edge filters
  - Directional filters + gating (non-linearity)
  - Sum/average pooling
  - Local response normalization (LRN)

see [Mahendran & Vedaldi, CVPR 2015]

HOG, dense SIFT, and many other “hand-engineered” features are convolutional feature maps.

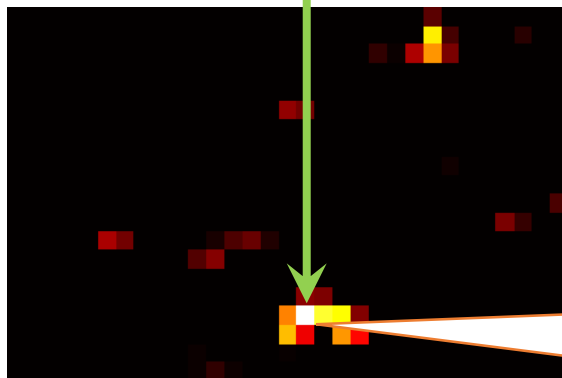


# Feature Maps = features and their locations



# Feature Maps = features and their locations

ImageNet images with **strongest** responses of this channel



one feature map of conv<sub>5</sub>  
(#55 in 256 channels of a model  
trained on ImageNet)



Intuition of *this* response:



There is a “**circle-shaped**” object (likely a tire) **at this position**.

**What**

**Where**

# Feature Maps = features and their locations

ImageNet images with **strongest** responses of this channel



one feature map of conv<sub>5</sub>  
(#66 in 256 channels of a model  
trained on ImageNet)



Intuition of *this* response:  
There is a “**λ-shaped**” object (likely an underarm) **at this position**.

**What**

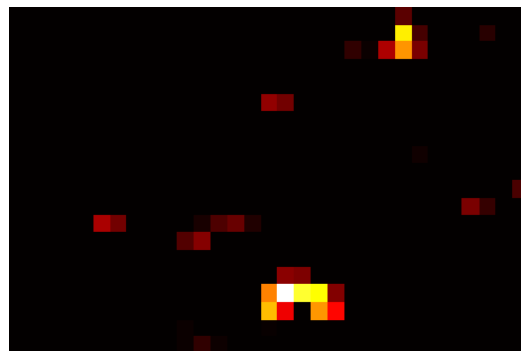
**Where**

# Feature Maps = features and their locations

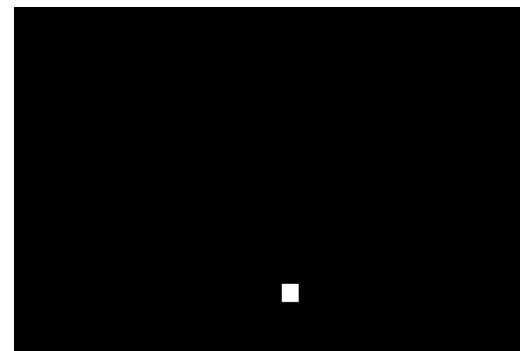
- Visualizing **one response** (by Zeiler and Fergus)



image



a feature map

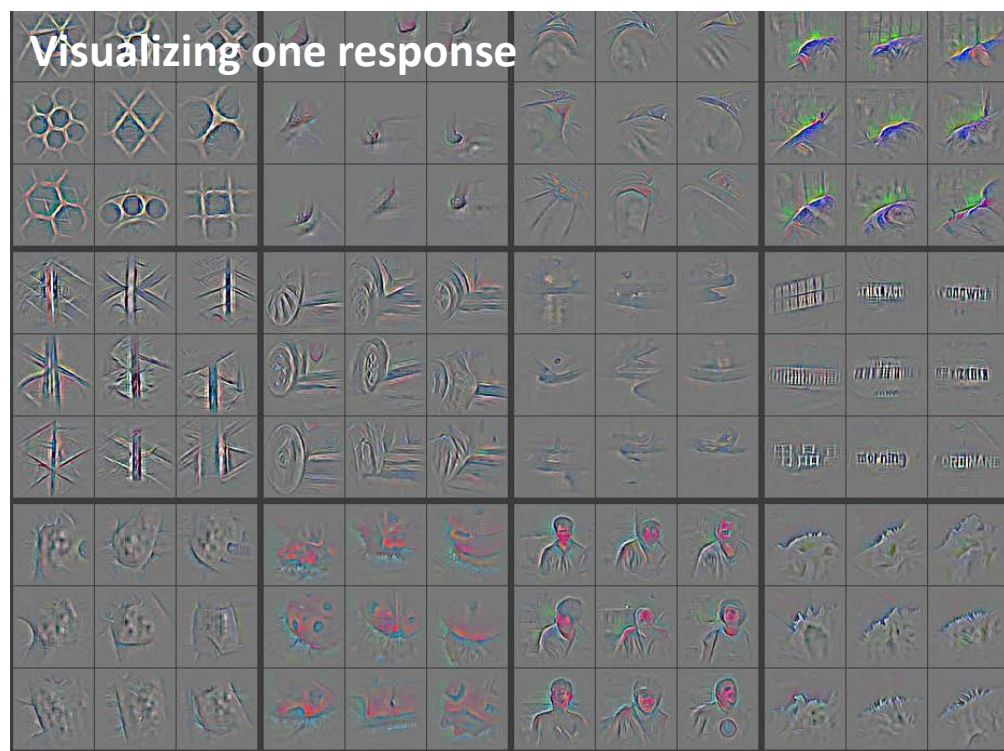


keep one response  
(e.g., the strongest)





# Feature Maps = features and their locations



conv3





image credit: Zeiler & Fergus



# Feature Maps = features and their locations




Intuition of *this* visualization:  
There is a “**dog-head**” shape **at this position**.

- **Location** of a feature:  
explicitly represents *where* it is.
- **Responses** of a feature:  
encode *what* it is, and implicitly  
encode finer position information –

*finer position information is  
encoded in the channel dimensions  
(e.g., bbox regression from  
responses at one pixel as in RPN)*


# Receptive Field



- Receptive field of the first layer is the filter size
- Receptive field (w.r.t. input image) of a deeper layer depends on all previous layers' filter size and strides

- **Correspondence** between a feature map pixel and an image pixel is not unique
- Map a feature map pixel to **the center of the receptive field** on the image in the SPP-net paper

# Receptive Field



## How to compute the center of the receptive field

- A simple solution
  - For each layer, pad  $\lfloor F/2 \rfloor$  pixels for a filter size  $F$  (e.g., pad 1 pixel for a filter size of 3)
  - On each feature map, the response at  $(0, 0)$  has a receptive field centered at  $(0, 0)$  on the image
  - On each feature map, the response at  $(x, y)$  has a receptive field centered at  $(Sx, Sy)$  on the image (stride  $S$ )

- A general solution

$$i_0 = g_L(i_L) = \alpha_L(i_L - 1) + \beta_L,$$

$$\alpha_L = \prod_{p=1}^L S_p,$$

$$\beta_L = 1 + \sum_{p=1}^L \left( \prod_{q=1}^{p-1} S_q \right) \left( \frac{F_p - 1}{2} - P_p \right)$$

See [Karel Lenc & Andrea Vedaldi]  
"R-CNN minus R". BMVC 2015.



# Region-based CNN Features

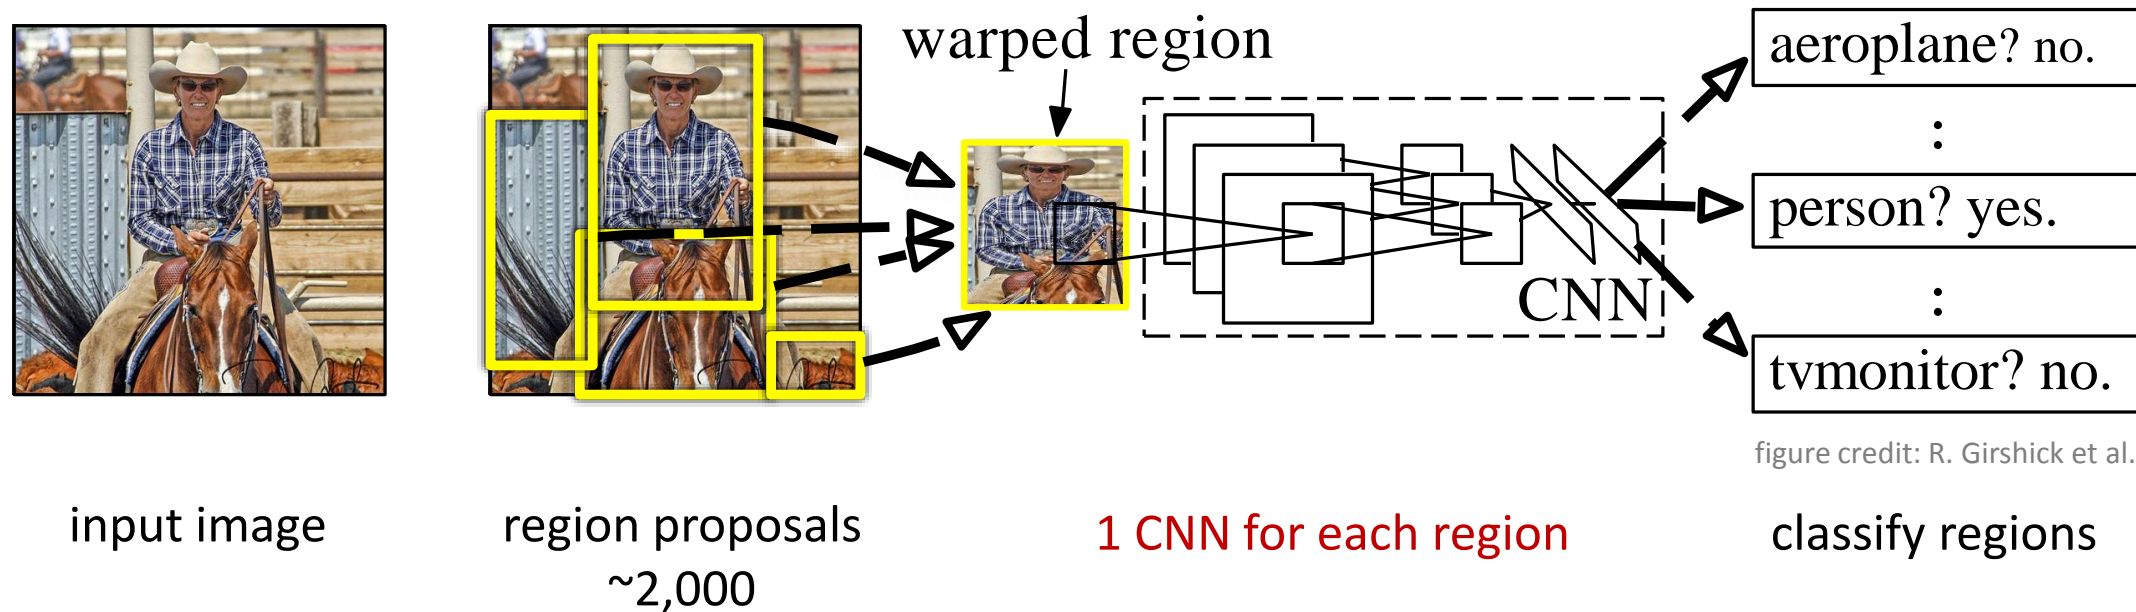




figure credit: R. Girshick et al.

## R-CNN pipeline

# Region-based CNN Features

- Given proposal regions, what we need is **a feature for each region**
- R-CNN: **cropping an image region** + CNN on region, requires 2000 CNN computations
- What about **cropping feature map regions**?



# Regions on Feature Maps




image region



feature map region

- Compute convolutional feature maps on the entire image **only once**.
- Project an image region to a **feature map region** (using correspondence of the receptive field center)
- Extract a region-based feature from the feature map region...

# Regions on Feature Maps



- **Fixed-length** features are required by fully-connected layers or SVM
- But how to produce a fixed-length feature from a feature map region?
- Solutions in traditional computer vision: Bag-of-words, SPM...

# Bag-of-words & Spatial Pyramid Matching

SIFT/HOG-based  
feature maps

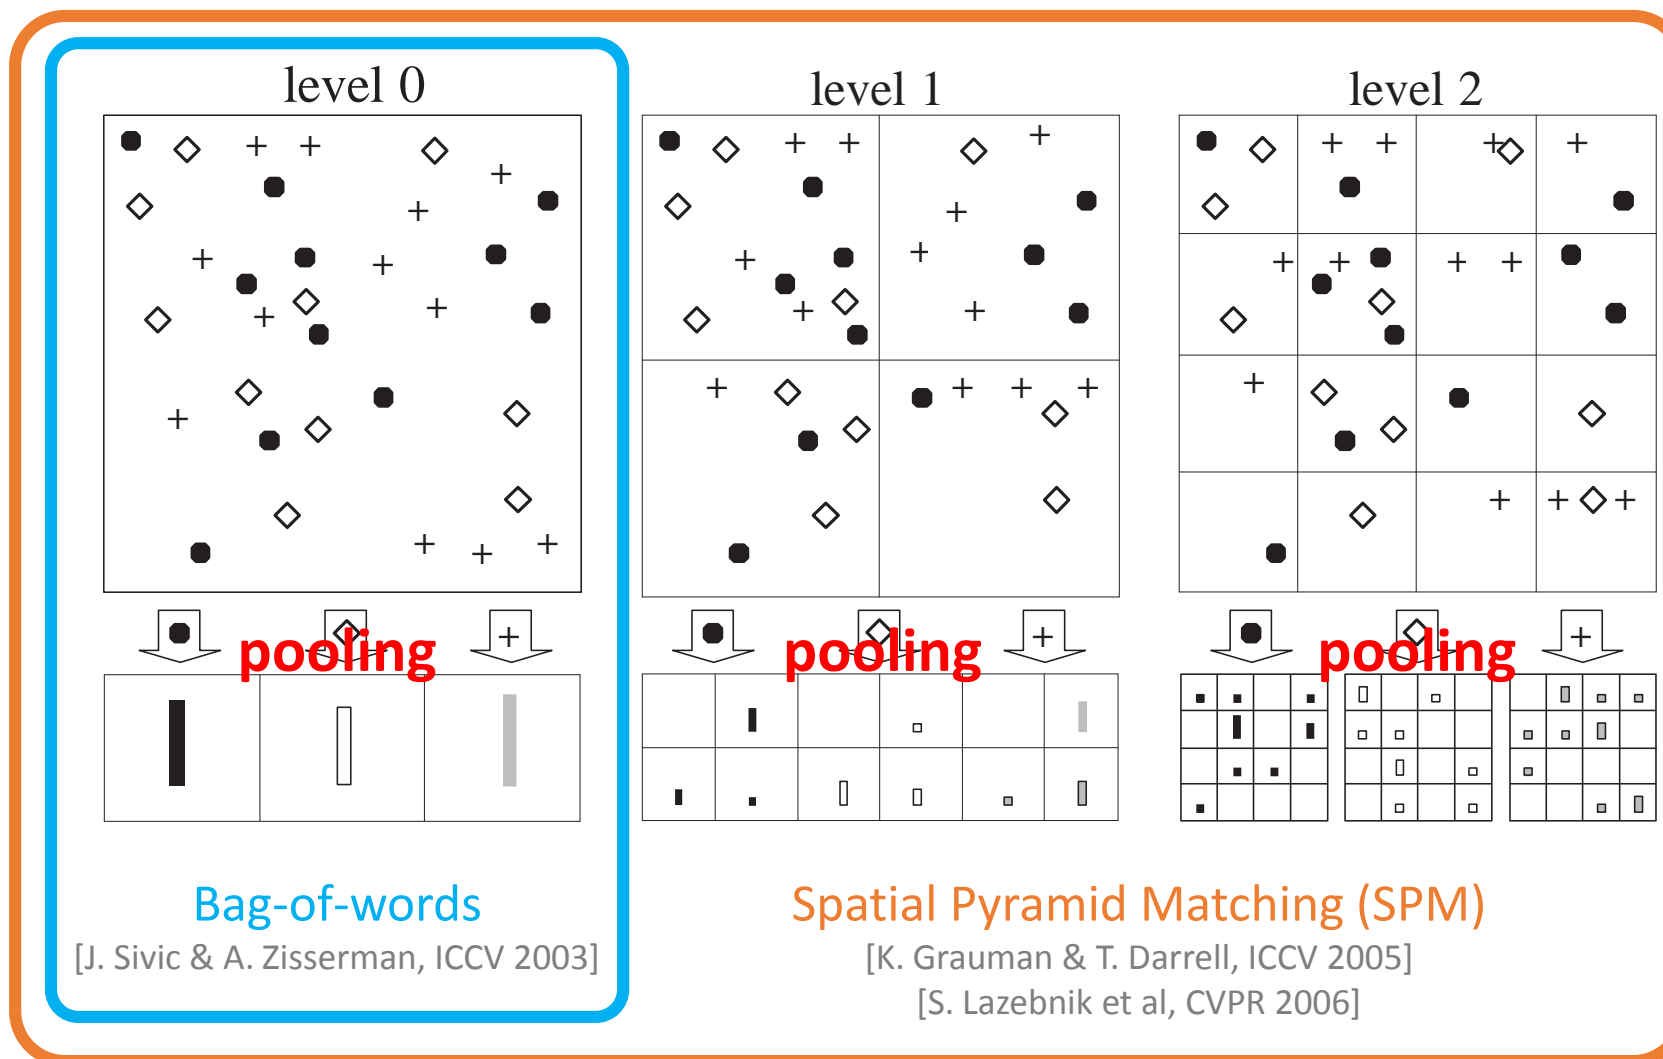





figure credit: S. Lazebnik et al.

# Spatial Pyramid Pooling (SPP) Layer

- fix the number of bins (instead of filter sizes)
- adaptively-sized** bins



**pooling**



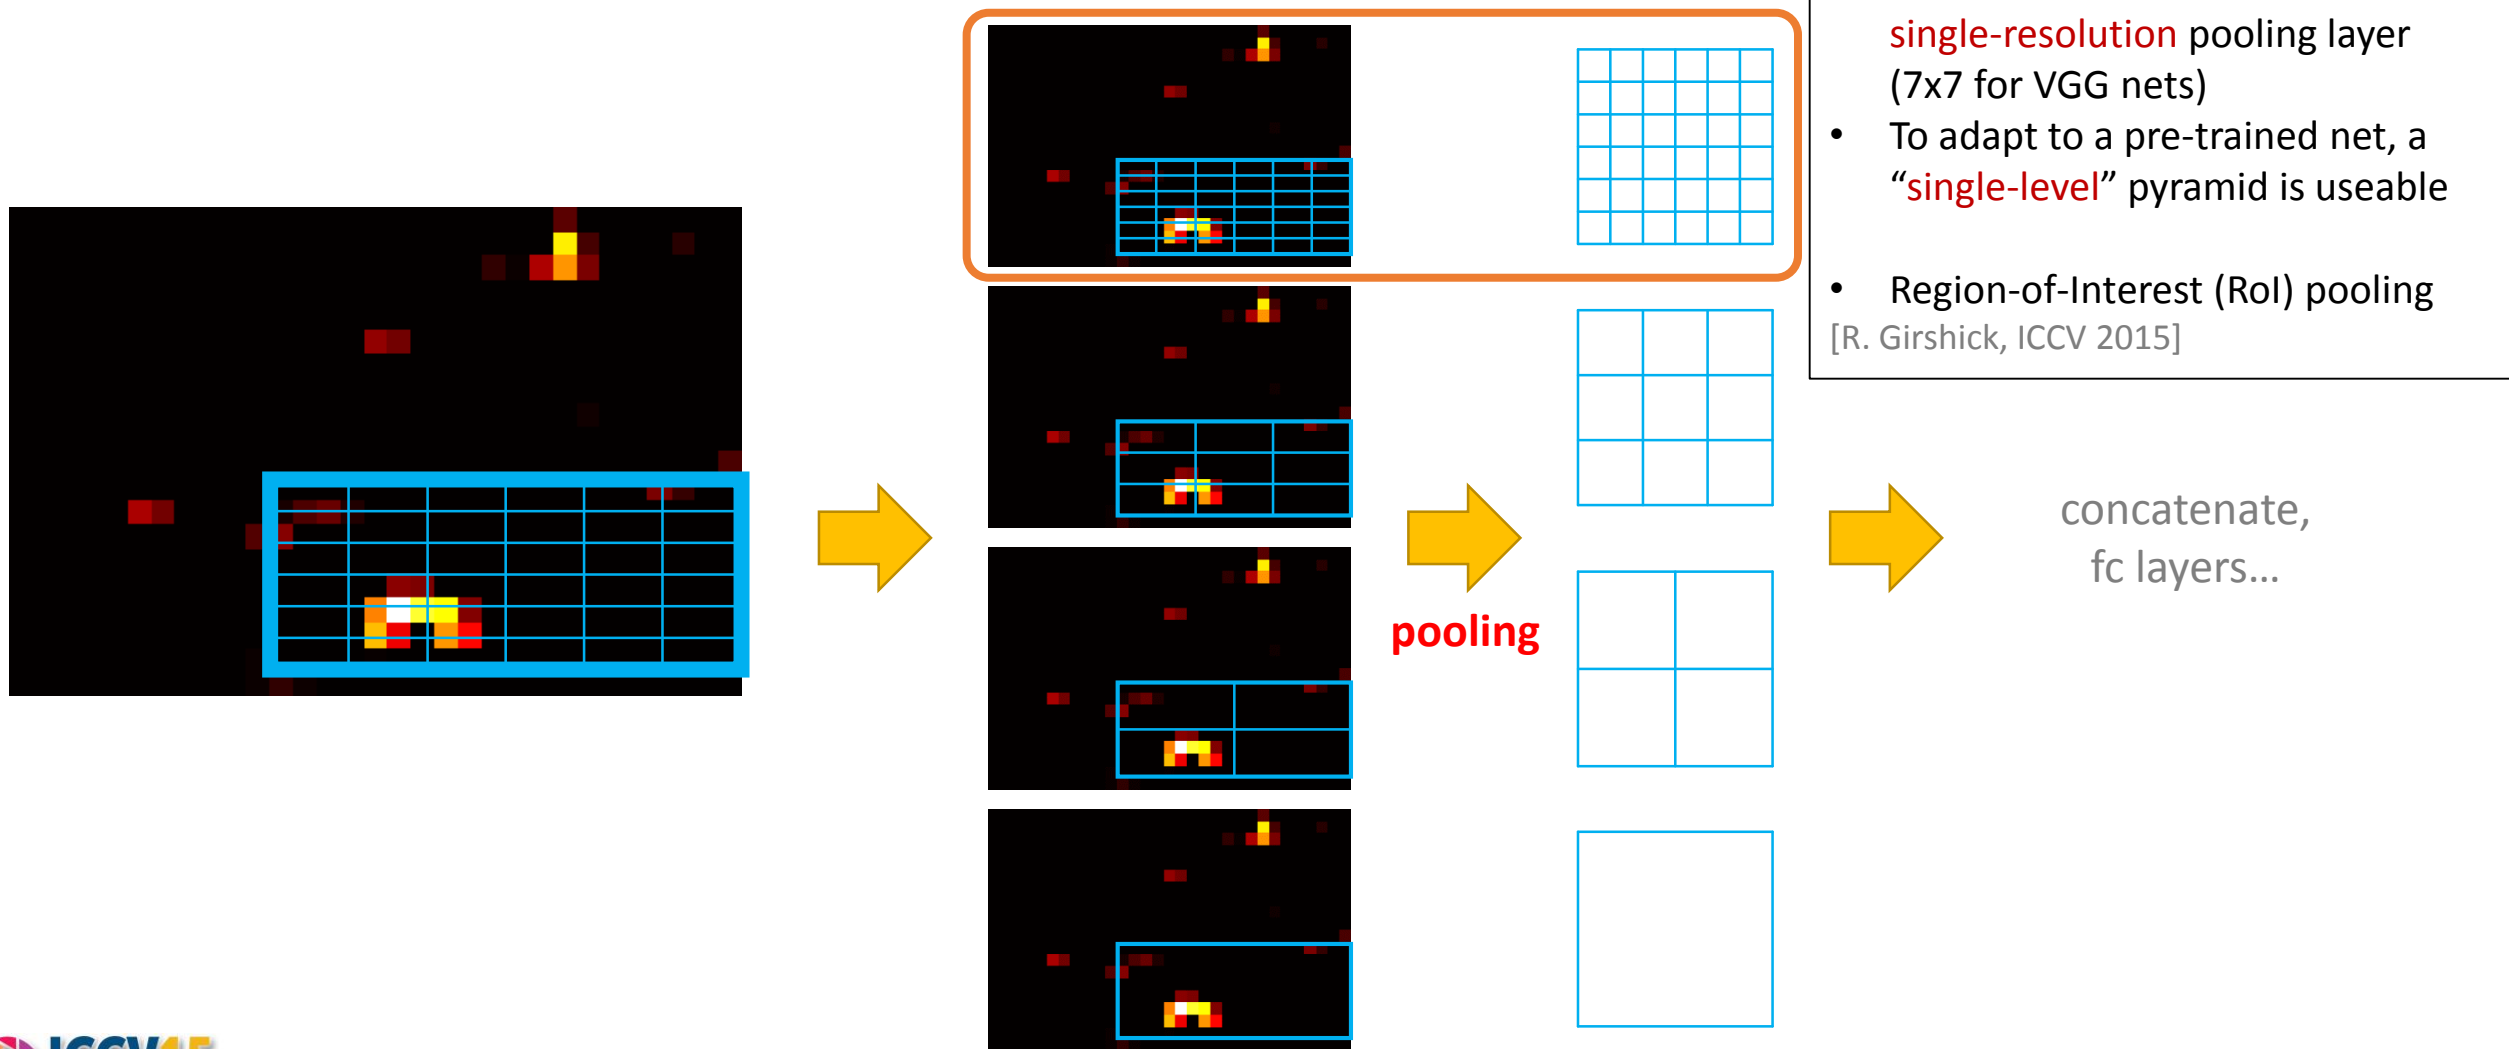
concatenate,  
fc layers...

a finer level maintains  
explicit spatial information

a coarser level removes  
explicit spatial information  
(bag-of-features)



# Spatial Pyramid Pooling (SPP) Layer



# Single-scale and Multi-scale Feature Maps

- Feature Pyramid
  - Resize the input image to multiple scales
  - Compute feature maps for each scale
  - Used for HOG/SIFT features and convolutional features (OverFeat [Sermanet et al. 2013])





# Single-scale and Multi-scale Feature Maps

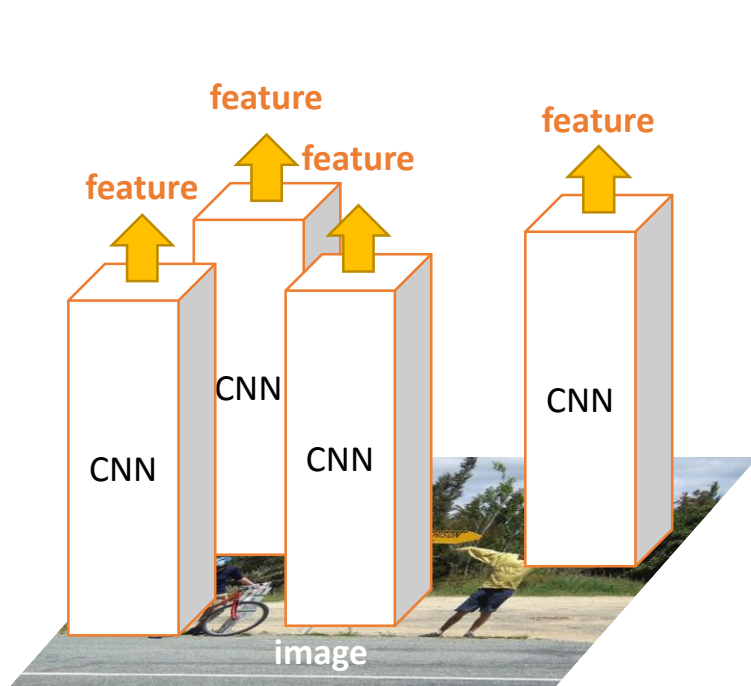
- But deep convolutional feature maps perform well **at a single scale**

	SPP-net <b>1-scale</b>	SPP-net 5-scale
pool <sub>5</sub>	43.0	44.9
fc <sub>6</sub>	42.5	44.8
fine-tuned fc <sub>6</sub>	52.3	53.7
fine-tuned fc <sub>7</sub>	54.5	55.2
fine-tuned fc <sub>7</sub> bbox reg	58.0	59.2
conv time	<b>0.053s</b>	0.293s
fc time	0.089s	0.089s
total time	0.142s	0.382s

- Also observed in Fast R-CNN and VGG nets
- Good speed-vs-accuracy tradeoff
- Learn to be scale-invariant from pre-training data (ImageNet)
- (note: but if good accuracy is desired, feature pyramids are still needed)


detection mAP on PASCAL VOC 2007, with ZF-net pre-trained on ImageNet  
this table is from [K. He, et al. 2014]

# R-CNN vs. Fast R-CNN (forward pipeline)



## R-CNN

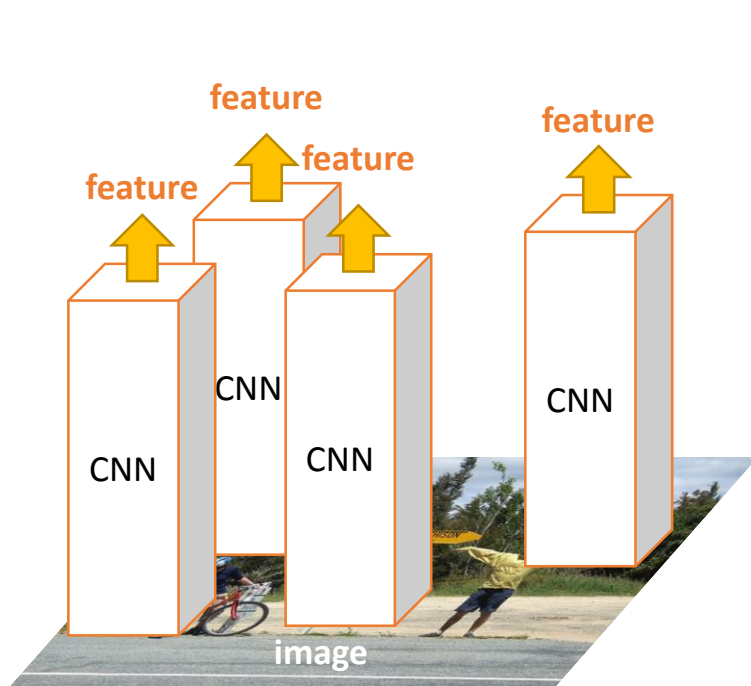
- Extract image regions
- 1 CNN per region (2000 CNNs)
- Classify region-based features



## SPP-net & Fast R-CNN (the same forward pipeline)


- 1 CNN on the entire image
- Extract features from feature map regions
- Classify region-based features

# R-CNN vs. Fast R-CNN (forward pipeline)



## R-CNN

- Complexity:  $\sim 224 \times 224 \times 2000$



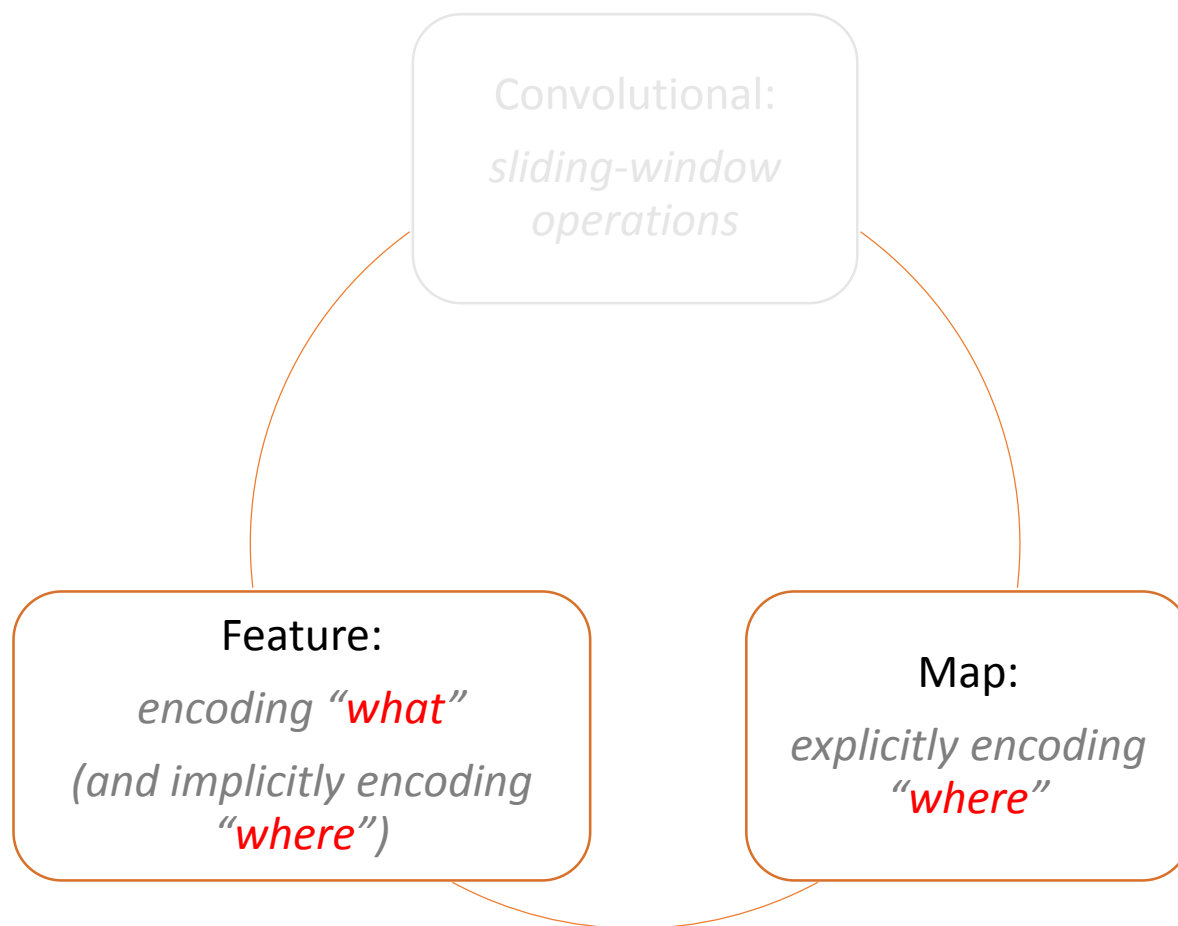
## SPP-net & Fast R-CNN (the same forward pipeline)

- Complexity:  $\sim 600 \times 1000 \times \mathbf{1}$
- $\sim \mathbf{160x}$  faster than R-CNN

# Region Proposal from Feature Maps

- Object detection networks are fast (0.2s)...
- but what about **region proposal**?
  - Selective Search [Uijlings et al. ICCV 2011]: 2s per image
  - EdgeBoxes [Zitnick & Dollar. ECCV 2014]: 0.2s per image
- Can we do region proposal **on the same set of feature maps**?

# Feature Maps = features and their locations

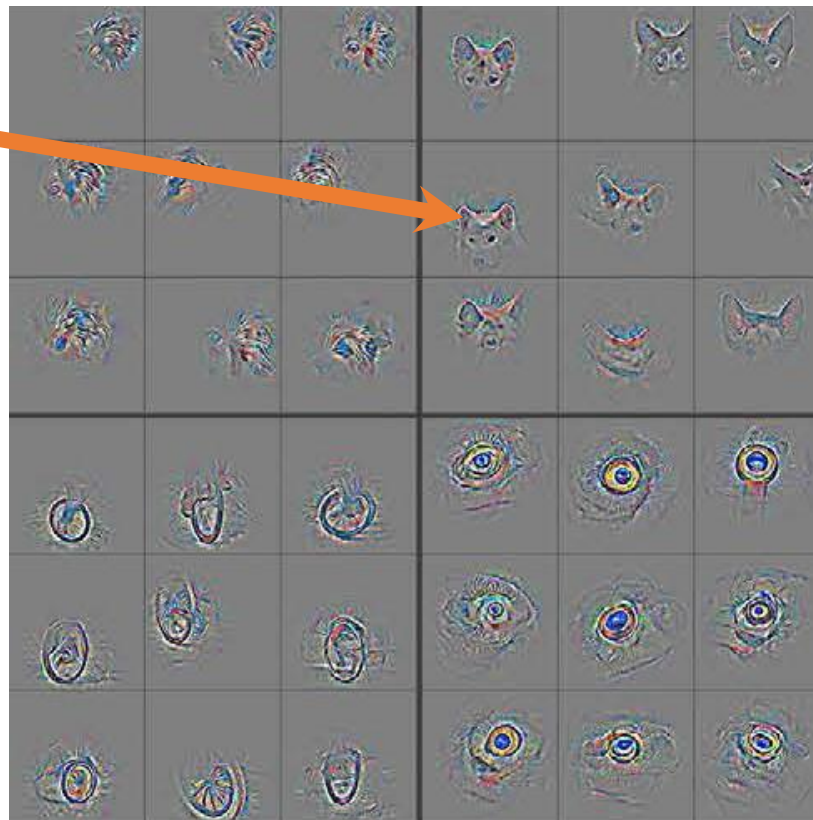


# Region Proposal from Feature Maps

- By decoding **one response** at a single pixel, we can still roughly see the object outline\*
- Finer localization information** has been encoded in the channels of a convolutional feature response
- Extract this information for better localization...


\* Zeiler & Fergus's method traces unpooling information so the visualization involves more than a single response. But other visualization methods reveal similar patterns.

Revisiting visualizations from Zeiler & Fergus



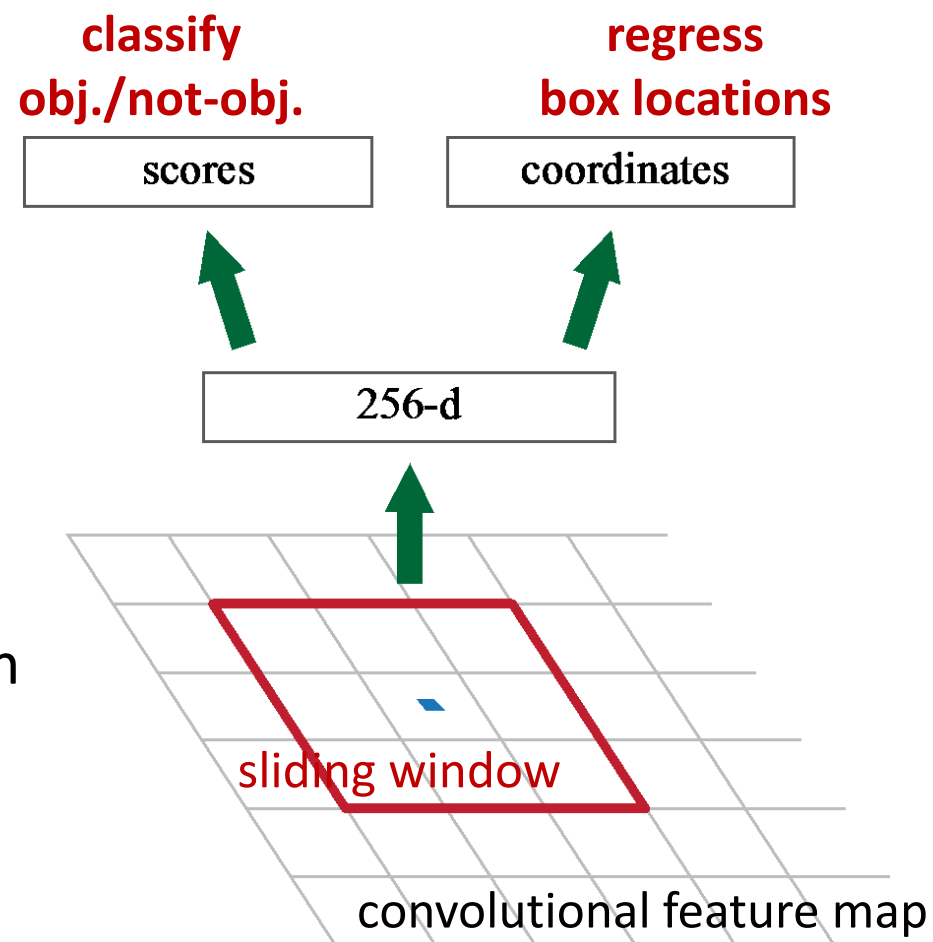


# Region Proposal from Feature Maps



# Region Proposal Network

- Slide a small window on the feature map
- Build a small network for:
  - classifying object or not-object, and
  - regressing bbox locations
- Position of the sliding window provides localization information **with reference to the image**
- Box regression provides finer localization information **with reference to this sliding window**






# Anchors as references

- **Anchors:** pre-defined reference boxes

- Box regression is with reference to anchors: regressing an anchor box to a ground-truth box


- Object probability is with reference to anchors, e.g.:

- anchors as positive samples: if  $\text{IoU} > 0.7$  or  $\text{IoU}$  is max
- anchors as negative samples: if  $\text{IoU} < 0.3$




# Anchors as references

- **Anchors**: pre-defined reference boxes
- **Translation-invariant** anchors:
  - the same set of anchors are used at each sliding position
  - the same prediction functions (with reference to the sliding window) are used
  - a translated object will have a translated prediction



# Anchors as references

- **Anchors**: pre-defined reference boxes
- **Multi-scale/size** anchors:
  - multiple anchors are used at each position:  
e.g., 3 scales ( $128^2$ ,  $256^2$ ,  $512^2$ ) and 3 aspect ratios (2:1, 1:1, 1:2) yield 9 anchors
  - each anchor has its own prediction function
  - **single-scale** features, multi-scale predictions



# Anchors as references


- Comparisons of **multi-scale** strategies



Image/Feature Pyramid




Filter Pyramid



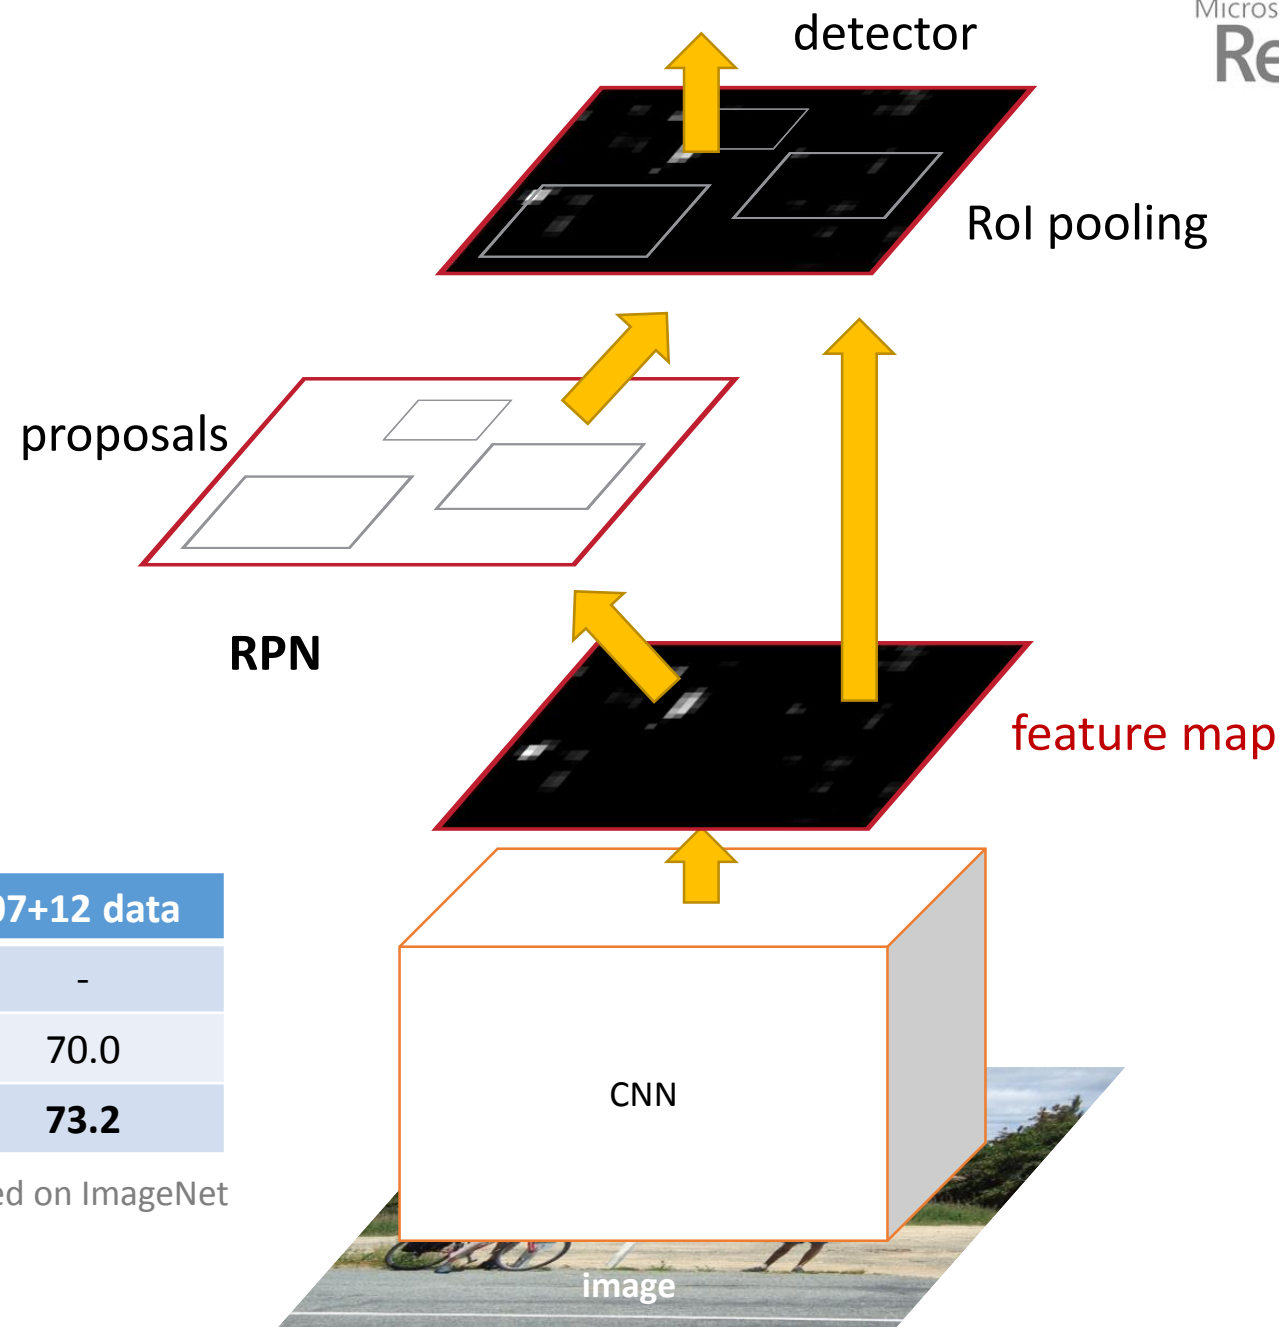
Anchor Pyramid

# Region Proposal Network

- RPN is **fully convolutional** [Long et al. 2015]
- RPN is trained end-to-end
- RPN **shares** convolutional feature maps with the detection network (covered in Ross's section)



# Faster R-CNN






system	time	07 data	07+12 data
R-CNN	~50s	66.0	-
Fast R-CNN	~2s	66.9	70.0
Faster R-CNN	<b>198ms</b>	<b>69.9</b>	<b>73.2</b>

detection mAP on PASCAL VOC 2007, with VGG-16 pre-trained on ImageNet



# Example detection results of Faster R-CNN



# Keys to efficient CNN-based object detection

- Feature **sharing**
  - R-CNN => SPP-net & Fast R-CNN: sharing features **among proposal regions**
  - Fast R-CNN => Faster R-CNN: sharing features **between proposal and detection**
  - All are done by shared **convolutional feature maps**
- Efficient multi-scale solutions
  - **Single-scale** convolutional feature maps are good trade-offs
  - **Multi-scale anchors** are fast and flexible



# Conclusion of this section

- Quick introduction to convolutional feature maps
  - Intuitions: into the “black boxes”
  - How object detection networks & region proposal networks are designed
  - Bridging the gap between “hand-engineered” and deep learning systems
- Focusing on forward propagation (inference)
  - Backward propagation (training) covered by Ross’s section